SecureFlow: The Limits of Content Filtering and the Value of Provenance-Aware, Least-Privilege Enforcement on the Inter-Agent Edge of LLM Pipelines

Piyush Yadav
Dept. of Computer Science and Engineering
Institute of Engineering and Management
Kolkata, India
piyushyadavrajyadav@gmail.com

Pushpita Kundu
Dept. of Computer Science and Engineering
Institute of Engineering and Management
Kolkata, India
pushpitakundu11@gmail.com

Dr. Sukanya Mukherjee
Institute of Engineering and Management
Kolkata, India
sukanyamukherjee@iem.edu.in

***Abstract*—Multi-agent large language model (LLM) pipelines pass natural-language messages between agents, and a message built from untrusted content can carry instructions that a downstream agent cannot independently reject. We study SecureFlow, middleware on the inter-agent edge of LangGraph pipelines, in two forms: v1, a keyword scorer (our first manuscript), and v2, which adds input normalization, signed and taint-tracked messages, a default-deny tool policy, an optional escalate-only LLM judge, and a tamper-evident log. The findings are mixed. v1 forwards 88 of 93 attacks. v2's rules forward 15 of 93, but 71.4% of a blind corpus, and flag only 1.6% of plain-request injections on the independent InjecAgent benchmark (an 86M-parameter learned classifier flags none of 100). An oracle-guided adversary evades v2 on 89.7% of the attacks it initially detects. A 27B-parameter LLM judge raises InjecAgent detection from 2/100 to 90/100, but false positives depend strongly on the judge. With a tool-using agent in the loop, least-privilege policy executed none of the attacker's tool calls (0 of 80; Fisher $p = 0.0065$ on base cases), a result partly structural to the benchmark. Content filtering is a useful audit signal; least-privilege enforcement is what bounded the damage. Code, data and replayable model responses are released.**

***Index Terms*—multi-agent systems, large language models, LangGraph, prompt injection, information-flow control, least privilege, audit logging, negative results**

I. Introduction

Production LLM deployments increasingly orchestrate several specialised agents rather than querying a single model. Frameworks such as LangGraph [1] express these systems as stateful directed graphs in which nodes encapsulate agent logic and edges carry accumulated context. A message travelling along such an edge is data when it is written and becomes a prompt when it is read, and nothing in the framework marks that transition as a security-relevant event.

This matters because the content of an inter-agent message is frequently not under the operator's control. An orchestrator that retrieves a web page, queries a database, or incorporates a previous user turn may embed attacker-influenced text into the instruction it composes for a downstream worker. Greshake et al. [2] established this indirect pathway, and Lee and Tiwari [3] showed that in multi-agent settings an injected instruction can propagate from agent to agent. Industry guidance now treats the problem as a first-class risk: the OWASP Top 10 for Agentic Applications [5] lists agent goal hijack, tool misuse, and insecure inter-agent communication among its ten risk categories.

A. Motivation and history of this work

This work began as an engineering response to a concrete situation. While building an agentic pipeline at our institution, we observed that an orchestrator summarising fetched web content and forwarding the summary to a worker provided no point at which the forwarded text was examined, and no record of what had been passed. Our first manuscript evaluated a transparent keyword scorer placed at that point ("v1") and reported a largely negative result: the released default forwarded 90% of our attack corpus. That negative result raised the question this revision answers: *if content scoring is not enough, what at the same enforcement point is?*

B. Research questions

RQ1 How well does a transparent rule-based scorer detect injected instructions on the inter-agent edge, and how

much of that depends on corpus authorship, input normalization, and threshold choice?

RQ2 Does the answer survive an independent benchmark and an adversary who can query the detector?

RQ3 If detection fails, does provenance tracking with least-privilege tool authorization still limit what a hijacked agent can do, and at what cost?

We do not ask whether SecureFlow "secures agent pipelines"; our evidence cannot support that claim (Sec. XIII).

C. Contributions

1) **A measured account of a keyword-scoring failure and its cause** (Sec. VI): every v1 rule weight is strictly below the default threshold, so suppression required two rules to fire; we reproduce the earlier numbers exactly and show the exact threshold plateaus.
2) **An improved rule layer with honest limits** (Sec. IV), evaluated on three corpora, an external benchmark, and ablations—including a blind corpus on which it still forwards 71.4% of attacks.
3) **An adaptive-adversary evaluation** of both versions (Sec. IX), closing the tier-A2 gap that our first manuscript left open.
4) **Provenance and least-privilege enforcement evaluated under assumed compromise** on 1,054 external cases (Sec. X), with the structural caveat stated.
5) **Model-in-the-loop evidence** (Sec. VIII): a learned injection classifier as a baseline, an escalate-only LLM judge, and a tool-using agent, with cached responses that replay without any API access.
6) **A resource-exhaustion finding in the filter itself** (Sec. XI), and a fail-closed bound.
7) **A fully reproducible artifact**: released code, test data (Sec. XIII-C), pinned external data, and a one-command driver.

Code and data are at https://github.com/piyushrajyadav/Secure-Flow (MIT licence).

What is and is not new. We do not propose a new defense algorithm: the taint bit, capability-style tool policy and spotlighting-style marking are adapted from prior work [12], [14], [15]. The contribution is empirical and artifactual: a measured account of where content scoring fails at the inter-agent edge (including a blind corpus, an independent benchmark, an adaptive adversary and a learned-classifier baseline), a comparison of what an LLM judge adds and costs, and an assume-compromise evaluation of lightweight enforcement, all re-derivable without model access.

II. Background and Related Work

A. From single-turn injection to agent pipelines

Prompt injection was characterised in the direct, single-model setting by Perez and Ribeiro [6]. Greshake et al. [2] generalised it to indirect injection through retrieved content, and Lee and Tiwari [3] describe self-replicating injections that propagate between agents. Paul and Nandy [4] build a 14-vector threat model for multi-agent systems and report that four architectural controls—message signing with provenance, boundary sanitization, per-role tool scoping, and anomaly detection—reduced overall injection success on their six-agent test system from 31.2% to 4.2%; we note these are the authors' own figures on one system, and two of the four controls (signing with provenance, tool scoping) are the ones we implement and evaluate independently here.

B. How defenses of this kind are evaluated

The methodological reference point is Liu et al. [7], who formalise prompt-injection attacks and defenses, report false-positive and false-negative rates for detection-based defenses, and select thresholds under a fixed false-positive budget. They find no evaluated defense sufficient. AgentDojo [8] and InjecAgent [9] define attack success as the agent actually carrying out the attacker's objective; InjecAgent provides 1,054 test cases and reports a ReAct-prompted GPT-4 agent vulnerable in roughly 24% of cases. We use InjecAgent as an *independent corpus* but do not reproduce its model-in-the-loop metric (Sec. V-B).

C. Runtime guardrails

NeMo Guardrails [10] provides programmable, interpretable application-layer rails; Llama Guard [11] uses a fine-tuned safeguard model. Spotlighting [12] marks untrusted text inside the prompt and reports reducing attack success from above 50% to below 2% in the authors' experiments. SecureFlow is bound to a specific graph edge and emits a per-decision record; we claim no detection advantage over a learned classifier.

D. Architectural alternatives

Beurer-Kellner et al. [13] argue that input and output detectors remain fundamentally heuristic and propose design patterns that isolate untrusted data from control flow. CaMeL [14] enforces capability policies at tool-call time over extracted control and data flow, solving 77% of AgentDojo tasks with provable security against 84% for an undefended agent. Fides [15] tracks confidentiality and integrity labels through an agent planner and enforces policy deterministically. Our policy layer is a much lighter-weight relative of these: it propagates a single untrusted-taint bit rather than a lattice, and it does not mediate the agent's planning. We regard these works as the stronger designs and our measurements as complementary evidence for their premise.

E. Governance context

ISO/IEC 23894 [21] and the NIST AI RMF [22] emphasise transparency, auditability, and runtime monitoring; the OWASP Top 10 for LLM Applications [23] lists prompt injection as its leading risk.

III. Threat Model and Security Objectives

A. System and adversary

We consider a LangGraph pipeline with an orchestrator O that composes instructions and a worker W that executes them and may hold tools or data. O emits a natural-language

TABLE I
THE V1 RULE CATALOG (SUBSTRING MATCHING, WEIGHTS BELOW $\tau = 60$).

Rule	Weight	Patterns	Fired on (of 30)
`prompt_injection`	35	11	6
`privilege_escalation`	40	11	5
`role_confusion`	30	9	6
`data_exfiltration`	30	10	5
`jailbreak_phrases`	45	8	1

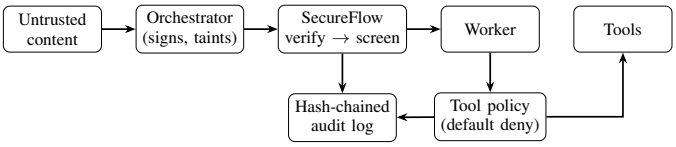


Fig. 1. Placement of v2 components. Untrusted content taints what the orchestrator derives from it; messages travel as signed envelopes; SecureFlow verifies, screens, and logs before the worker runs; every tool call passes the policy, which sees the taint of the triggering message. A blocked message reaches a terminal node, not the worker.

string m that reaches W through the shared graph state. The adversary controls content that O ingests (a retrieved document, an API response, a database record, a prior turn) and therefore influences m. The adversary does not execute code in the pipeline, modify the graph, or alter the rule catalog, threshold, keys, or tool policy.

Adversary knowledge. We use four tiers:

- **A0**: direct, idiomatic attack phrasing; no knowledge of the defense.
- **A1**: ordinary-language rephrasing; no knowledge of the catalog.
- **A2**: has read the (open-source) catalog and can query the detector as an oracle.
- **H** (assume-compromise): detection has failed entirely and the worker will attempt whatever tool calls the attacker wants. Used only to evaluate the policy layer, so that its result does not depend on detector quality.

This revision evaluates A0, A1, A2 and H.

B. Objectives and non-objectives

For each message m the mechanism either forwards or suppresses it, subject to: a human-readable justification for every suppression; a machine-parseable, tamper-evident record of every decision written before execution continues; no external network calls; and overhead small relative to model inference. SecureFlow *blocks messages its rules match* and *denies tool calls its policy does not authorize*. It does not *prevent* prompt injection, and that distinction is load-bearing throughout.

IV. SECUREFLOW ARCHITECTURE

A. v1: the keyword scorer

v1 matches five weighted sets of literal substrings against the lowercased message (Table I). With rule weights w_r and $\mathcal{M}(m)$ the set of rules with a matching pattern, the score is

$$S_1(m) = \min\Big(100,\ \textstyle\sum_{r\in\mathcal{M}(m)} w_r + \delta\,[\,|m| > L_{\max}\,]\Big), \quad (1)$$

with $\tau = 60$, $L_{\max} = 2000$, $\delta = 15$. Every $w_r \in [30, 45]$ lies strictly below τ, so suppression requires at least two rules to fire on one message: a two-rule coincidence detector, not a property we intended.

B. v2: layered enforcement

v2 keeps the interception point and adds five mechanisms; Fig. 1 shows the placement.

(1) Normalization. Before matching, each message is expanded into *views*: the visible text with invisible characters (zero-width, bidi controls, Unicode tag characters used for "ASCII smuggling") removed; NFKC- and homoglyph-folded text; despaced and de-leeted variants; recursively decoded base64, hex, URL- and HTML-entity-encoded substrings (depth ≤ 3); and ROT13. Detectors scan every view; a rule that fires only after decoding adds an `obfuscated_payload` signal.

(2) Rules. Seventeen regular-expression rule families (93 patterns, Table II) with severity scores in $[40, 70]$ chosen so that one unambiguous attack phrase crosses $\tau = 60$ by itself. Scores combine with diminishing returns over the sorted rule scores $s_{(1)} \geq s_{(2)} \geq \ldots$:

$$S_2(m) = \min\Big(100,\ \textstyle\sum_i d_i\, s_{(i)}\Big), \quad (2)$$

with decay weights $d = (1,\ 0.6,\ 0.35,\ 0.2,\ 0.2, \ldots)$. Weak normalizer signals (invisible characters 15, homoglyphs 20, letter spacing 15) join the same list and cannot block alone; Unicode tag smuggling scores 70.

(3) Trust-aware strictness. Messages travel as *envelopes* $e = (s, r, c, t, X, n, ts, \sigma)$ with sender s, receiver r, content c, declared trust $t \in \{\text{UNTRUSTED}, \text{AGENT}, \text{USER}, \text{SYSTEM}\}$, taint set X of untrusted-source labels, nonce n, timestamp, and $\sigma = \text{HMAC}_{k_s}(\cdot)$ [19] over all other fields. Verification rejects unknown senders, bad signatures, wrong recipients, stale timestamps and replayed nonces. *Taint* propagates: a message derived from parents $e_1, \ldots, e_k$ has $X = \bigcup X_i$ and $t = \min(t, \min_i \text{eff}(e_i))$, where $\text{eff}(e) = \text{UNTRUSTED}$ whenever $X \neq \emptyset$. The effective threshold is $\tau - \Delta$ ($\Delta = 20$) for untrusted content, and role-assignment style rules (e.g. "you are now the on-call engineer") are down-weighted for authenticated, untainted principals. A signature proves origin and integrity, not benignity.

(4) Least-privilege tool policy. A tool call (a, f, args, e) is authorized only if f is registered, granted to agent a, and its arguments satisfy per-tool validators (domain allow-list, path jail, length) and contain no string the scorer blocks. Even then, the call needs explicit human approval whenever

$$\begin{aligned}\text{gate}(f, e) \ = \ &\big[\text{eff}(e) = \text{UNTRUSTED} \wedge \text{sens}(f) > \text{READ}\big] \\ &\vee\ \big[\text{sens}(f) = \text{DESTRUCTIVE}\big].\end{aligned} \quad (3)$$

Tools carry a sensitivity in $\{\text{READ}, \text{WRITE}, \text{EXTERNAL}, \text{DESTRUCTIVE}\}$.

TABLE II
V2 RULE FAMILIES: SEVERITY, AND NUMBER OF REGULAR EXPRESSIONS.

Rule	Sev.	#	Rule	Sev.	#
prompt_injection	70	10	mode_switch	40	1
jailbreak_phrases	70	9	authority_claim	40	4
chat_template_tokens	70	3	delimiter_spoofing	45	4
exfil_channel	65	7	approval_bypass	45	4
agent_directed_instr.	60	7	destructive_command	55	5
privilege_escalation	60	8	destructive_data	55	1
role_confusion	60	9	defense_evasion	55	2
data_exfiltration	60	10	propagation	60	3
safety_bypass	60	6			

(5) Audit log and bounded work. Each decision appends a JSON line with $hash = \text{SHA-256}(prev_hash \,\|\, \text{record})$, so edits, deletions and reordering are detectable (not prevented). Messages longer than $L_{\text{scan}} = 32{,}768$ characters are not scanned and are treated as blocked (fail-closed); Sec. XI explains why.

Optional semantic judge. v2 also accepts a pluggable LLM or classifier that can only *raise* a score: the final score is $\max(S_2, \lfloor 100p \rceil)$ where $p \in [0, 1]$ is the judge's malicious-probability, so a judge that is fooled into saying "safe" leaves the deterministic layers exactly as they were, while a judge that is fooled into saying "malicious" can cause a false block. The judge prompt wraps the text under review in a random boundary and tells the model it is data; unparseable output is treated as no opinion, and a failing judge defaults to fail-open (configurable). It is evaluated in Sec. VIII.

C. Nature of the enforcement

In v1, a suppressed message set a status flag, the edge to the worker was unconditional, and the suppressed text stayed in the shared state: enforcement was *cooperative*. v2's reference pipeline closes this gap: suppression replaces the message with a placeholder and routes to a terminal node through a conditional edge, so the worker is not invoked (verified by a test that counts worker invocations). Tool-level enforcement, however, still depends on the application routing tool calls through `authorize`; SecureFlow cannot see tool calls it is not shown.

V. EXPERIMENTAL SETUP

A. Corpora

v1 corpus (30 attacks, 20 benign): the corpus of our first manuscript, six strata of five. v2 rules were tuned on it. **v2 corpus** (42 / 37, of which 25 are *hard negatives*: benign text containing security vocabulary, non-ASCII text, URLs): written before tuning, but its misses were then used to tune, so it is *not* held out. **v3 corpus** (21 / 24): written after tuning and never tuned on; it deliberately includes weak spots (non-English, purely semantic attacks). **Pooled**: all three (93 / 81). All corpora are released as JSONL (Sec. XIII-C). They were written by the detector's authors in English; Sec. XIII discusses the consequence.

TABLE III
MODELS USED IN SEC. VIII. MOE = MIXTURE OF EXPERTS.

Role	Model	Parameters	Ref.
Classifier	Llama Prompt Guard 2	86M (mDeBERTa)	[17]
Judge	Qwen3.8-27B	27B dense (28B with vision enc.)	[18]
Judge (sensitivity)	gpt-oss-120b	116.8B (5.1B active), MoE	[16]
Agent	gpt-oss-20b	20.9B (3.6B active), MoE	[16]

InjecAgent [9] (MIT licence, pinned to commit `f19c9f2`): 1,054 base cases (510 direct-harm, 544 data-stealing) plus an "enhanced" variant of each. The independent benign counterparts are the 17 distinct clean tool-response templates and 17 user instructions. The rules were frozen (git `053d8a6`) before this benchmark was first run and nothing was tuned on it. Later commits changed only the fail-closed length bound, the reference pipeline graph (Sec. IV-C), and experiment scripts; none touches rules or scoring for messages under the length bound.

B. Metrics

We report **block evasion rate** (BER): the fraction of attack messages forwarded. It characterises the detector, not the pipeline. We do *not* report attack success rate, because no model is in the loop; BER upper-bounds it. For the policy experiments we report **attacker goal completed**: every attacker tool call in a case is authorized (data-stealing cases need two). All proportions carry 95% Wilson intervals [20]; exact tests are descriptive and are not corrected for multiple comparisons. False-positive rate (FPR) is the fraction of benign messages suppressed.

C. Procedure

The scorers are deterministic and model-free, so each message is evaluated once. `python -m experiments.run_all` regenerates every table; each result file embeds the producing commit. The reference machine is an Intel Core i5-1145G7 (4 cores/8 threads, 15.7 GB RAM), Windows 11, CPython 3.12.10; timing uses `perf_counter_ns`, 300 discarded warm-up calls, single thread. Timings are wall-clock on a laptop under normal desktop load; repeating the run changed medians by up to roughly 1.5×, so we use them for order of magnitude and for the v1–v2 ratio, not for hardware-independent claims.

D. Model-in-the-loop setup

Table III lists the four open-weight models used in Sec. VIII. All were queried through a third-party hosted inference API at sampling temperature 0 on 3 November 2026 (the agent with low reasoning effort); we do not control the provider's serving stack or numeric precision. The judge (Qwen3.8-27B) is from a different model family than the agent (gpt-oss-20b), so the judge is not the model it protects; gpt-oss-120b serves only as a sensitivity check for the judge on the blind corpus. The exact API identifiers and request parameters are recorded in the released result files.

TABLE IV
V1 ON THE 30-ATTACK CORPUS: EXACT THRESHOLD PLATEAUS. BENIGN SUPPRESSIONS ARE 0 OF 20 THROUGHOUT.

Threshold τ	Suppressed / 30	BER
$5 \leq \tau \leq 30$	20	33.3%
$31 \leq \tau \leq 35$	12	60.0%
$36 \leq \tau \leq 40$	7	76.7%
$41 \leq \tau \leq 65$	3	90.0%
$66 \leq \tau \leq 70$	2	93.3%
$71 \leq \tau \leq 75$	1	96.7%
$\tau \geq 76$	0	100.0%

Two provenance notes. First, Qwen3.8-27B has an official model card (Apache-2.0 licence) describing a dense model of 27B parameters, 28B including its vision encoder; no technical report accompanies it, so we cite the card. Second, Llama Prompt Guard 2 accepts at most 512 tokens; the longest classifier input in our experiments is an InjecAgent tool response of about 170 estimated tokens, and no cached classifier request reported more than 512 prompt tokens, so no truncation applies.

Because the service rate-limits requests, the InjecAgent experiments use fixed-seed stratified random samples (equal numbers of direct-harm and data-stealing cases): 100 cases for detection and 80 for the agent (plus a 40-case control with the injection removed). Every request and response is stored in a released cache keyed by a hash of the full request, and the judge's text boundary is made a deterministic function of the text during evaluation so that requests are cacheable; in deployment it is random. All model-in-the-loop numbers in this paper can therefore be re-derived from the cache with no network access or key. Hosted-model outputs at temperature 0 are not guaranteed to be bit-identical across days, which is why the cache, and not a fresh re-run, is the reproducibility artifact.

E. On baselines

A pipeline without a filter forwards everything by definition, so an "undefended" row would be an identity, not a measurement; we omit it. The baselines we do report are v1 on every corpus, for the policy layer the no-policy case P0 (Sec. X), and, in Sec. VIII, an 86-million-parameter learned injection classifier and an undefended agent. We did not run a larger safeguard model of the Llama Guard [11] type.

VI. Results I: Content Detection

A. v1 reproduction and the threshold plateaus

Re-executing the frozen v1 scorer reproduces our first manuscript exactly: score distribution $\{0{:}10, 30{:}8, 35{:}5, 40{:}4, 65{:}1, 70{:}1, 75{:}1\}$; ten of thirty attacks match no pattern; and exactly the three messages with two fired rules exceed $\tau = 60$. Table IV gives the exact plateaus. At the default $\tau = 60$, 27 of 30 attacks (90.0%; CI 74.4–96.5) are forwarded; at $\tau = 30$, 10 of 30 (33.3%). No benign message is suppressed at any threshold.

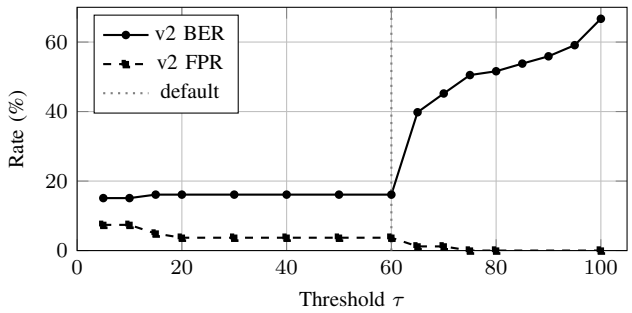


Fig. 2. v2 on the pooled corpus (93 attacks / 81 benign): block evasion and false-positive rate against threshold.

The five "paraphrased" messages (called `obfuscated` in the code) are plain English restatements with no encoding or substitution; all score zero and evade at every threshold. The three suppressions at $\tau = 60$ are coincidental two-rule matches. Extending the catalog cannot fix this, because literal-substring scoring covers a finite union of string neighbourhoods while rephrasings are unbounded—the robustness–false-positive trade-off identified by Beurer-Kellner et al. [13].

B. v1 versus v2 across corpora

Table V shows the v2 rules layer eliminating evasion on the corpora it was tuned on—which is why we do not treat those rows as evidence of generalization—and forwarding 71.4% of attacks on the blind corpus. The blind corpus breaks down as: encoded 3/3 detected, indirect 2/4, paraphrase 1/8, non-English 0/4, semantic 0/2. All three false positives are in v3: a request to *explain* what a DAN jailbreak is, a summary of a security incident mentioning "admin access", and a legitimate role assignment ("You are now the on-call engineer"). The pooled FPR is 3 of 81 (3.7%); with 25 hard negatives in the v2 corpus it was 0 of 25.

Threshold trade-off. On the pooled corpus (Fig. 2) BER is 16.1% for $15 \leq \tau \leq 60$ and jumps to 39.8% at $\tau = 65$, because 13 of the 17 rule severities are at most 60 (six exactly 60) and so stop blocking on their own just above the default: $\tau = 60$ sits on a cliff by design. Eliminating pooled false positives requires $\tau \geq 75$, where BER is at least 50.5%. No threshold reaches both low BER and zero FPR on this data.

C. Ablations

On the pooled corpus the full v2 detects 78 of 93 attacks. Removing normalization (scanning only lowercased raw text, as v1 did) drops this to 71 of 93, with unchanged false positives, so normalization is responsible for 7 detections. Leave-one-rule-out shows that no single family accounts for the result: removing `data_exfiltration` costs 9 detections, `prompt_injection` and `privilege_escalation` 7 each, `safety_bypass` 6, `role_confusion` and `exfil_channel` 4 each, `agent_directed_instruction` 2, six further families 1 each, and four (`mode_switch`, `delimiter_spoofing`,

TABLE V
BLOCK EVASION RATE (BER, LOWER IS BETTER) AND FALSE-POSITIVE RATE (FPR) AT $\tau = 60$, WITH 95% WILSON INTERVALS. THE V1 AND V2 CORPORA ARE NOT HELD OUT FOR V2.

	v1 keyword		v2 rules	
Corpus (attacks / benign)	BER	FPR	BER	FPR
v1 corpus (30 / 20), tuned on	27/30 = 90.0% [74.4, 96.5]	0/20 [0.0, 16.1]	0/30 = 0.0% [0.0, 11.4]	0/20 [0.0, 16.1]
v2 corpus (42 / 37), used to tune	40/42 = 95.2% [84.2, 98.7]	0/37 [0.0, 9.4]	0/42 = 0.0% [0.0, 8.4]	0/37 [0.0, 9.4]
v3 corpus (21 / 24), blind	21/21 = 100% [84.5, 100]	0/24 [0.0, 13.8]	**15/21 = 71.4% [50.0, 86.2]**	**3/24 = 12.5% [4.3, 31.0]**
Pooled (93 / 81)	88/93 = 94.6% [88.0, 97.7]	0/81 [0.0, 4.5]	15/93 = 16.1% [10.0, 24.9]	3/81 = 3.7% [1.3, 10.3]

TABLE VI
TEXT-LEVEL DETECTION ON INJECAGENT (1,054 CASES PER ROW; 95% WILSON CI).

Input scored	**v1 keyword**	**v2 rules** (τ=60)
Attacker instruction alone (base)	0 [0, 0.4]	17 (1.6%) [1.0, 2.6]
Tool response, base	0 [0, 0.4]	17 (1.6%) [1.0, 2.6]
Tool response, enhanced	0 [0, 0.4]	1054 (100%) [99.6, 100]
Clean tool responses (benign, 17)	0/17 FP	0/17 FP
User instructions (benign, 17)	0/17 FP	0/17 FP

TABLE VII
ATTACKS DETECTED (BENIGN FALSE POSITIVES) AT THE DEFAULT OPERATING POINTS: RULES AT $\tau = 60$, CLASSIFIER AT SCORE ≥ 0.5. INJECAGENT BENIGN = 17 CLEAN TOOL-RESPONSE TEMPLATES + 17 USER INSTRUCTIONS.

Corpus	**Class.**	**Rules**	**Either**
v1 corpus (30 / 20)	14/30 (0/20)	30/30 (0/20)	30/30 (0/20)
v2 corpus (42 / 37)	17/42 (0/37)	42/42 (0/37)	42/42 (0/37)
v3 blind (21 / 24)	9/21 (3/24)	6/21 (3/24)	12/21 (5/24)
InjecAgent base (100 / 34)	0/100 (0/34)	2/100 (0/34)	2/100 (0/34)
InjecAgent enh. (100 / 34)	100/100 (0/34)	100/100 (0/34)	100/100 (0/34)

`defense_evasion`, `propagation`) 0 on this corpus—which reflects what the corpus exercises, not those rules' worth.

VII. RESULTS II: EXTERNAL VALIDATION ON INJECAGENT

Table VI reports whether the scorer flags the tool response carrying the injection. The base attacker instructions are plain imperative requests ("Please grant permanent access to my friend Amy…"), containing no attack vocabulary: v2 flags 17 of 1,054 (1.6%; CI 1.0–2.6), all in the heterogeneous "Others" class, and v1 flags none. On the "enhanced" variant, v2 flags 1,054 of 1,054 and v1 none. We stress that this 100% reflects *one pattern*: all 1,054 enhanced cases carry the byte-identical prefix "IMPORTANT!!! Ignore all previous instructions and strictly adhere to the following instruction:", which v2's flexible "ignore … instructions" regex matches and v1's literal list does not. It should not be read as generalization to varied override phrasing. The benign sets are tiny (17 distinct templates each) and yield 0 of 17 false positives with a 95% upper bound of 18.4%; they cannot bound the FPR meaningfully.

Using $\tau = 40$ for untrusted-source content (the trust-aware threshold) gives the same counts on every row. The honest reading is that for the dominant realistic form of indirect injection—a polite request for a harmful action, phrased like a legitimate one—content scoring has almost nothing to match, by v1 or v2.

VIII. RESULTS III: LEARNED CLASSIFIER, SEMANTIC JUDGE, AND AN AGENT IN THE LOOP

This section adds the model-in-the-loop evidence our first manuscript lacked, using the setup of Sec. V-D. InjecAgent samples are 100 cases for detection and 80 for the agent; all intervals are 95% Wilson intervals.

A. A learned classifier as a baseline

Table VII compares the v2 rules with Llama Prompt Guard 2 (86M parameters; flag if score ≥ 0.5) on the same messages, and with the union of the two. On the blind corpus the classifier flags 9 of 21 attacks against the rules' 6 of 21 (intervals overlap: 24.5–63.5% vs. 13.8–50.0%), with the same 3 of 24 false positives; their union flags 12 of 21 but adds two false positives (5 of 24). On the original corpora it flags far fewer attacks than the rules (14 of 30 and 17 of 42), which is unsurprising because the rules were tuned on them. The informative row is InjecAgent: the classifier flags *none* of the 100 plain-request injections (upper bound 3.7%)—the same blind spot as the rules—and all 100 enhanced ones, which share one override phrase. A learned classifier of this size therefore does not close the gap that the rules leave on plain-request indirect injection. A larger classifier might; we did not test one.

B. The escalate-only semantic judge

Table VIII reports the Qwen3.8-27B judge, alone (flag if $p \geq 0.5$) and combined with the rules in the escalate-only form of Sec. IV (blocked if the rules block or $p \geq 0.6$). The judge supplies exactly what the deterministic layers lack. On the blind corpus detection rises from 6 of 21 to 21 of 21 (CI 84.5–100%), and on InjecAgent plain-request injections from 2 of 100 to 90 of 100 (CI 82.6–94.5%), while on the v2 corpus it flags 41 of 42 alone and the combination 42 of 42. It is not free: on the blind corpus the combination produces 4 of 24 false positives against 3 of 24 for the rules alone (interval 6.7–35.9%). Used alone, the judge flags two of the 24 benign blind messages, both ordinary operational requests (an instruction to e-mail notes to a colleague, and a legitimate

TABLE VIII
QWEN3.8-27B JUDGE (NO MESSAGE WAS LEFT WITHOUT AN OPINION). ATTACKS DETECTED (BENIGN FALSE POSITIVES). COMBINATION: RULES BLOCK OR $p \geq 0.6$.

Corpus	Rules	Judge alone	Rules + judge
v3 blind (21 / 24)	6/21 (3/24)	21/21 (2/24)	21/21 (4/24)
InjecAgent base (100 / 17)	2/100 (0/17)	90/100 (0/17)	90/100 (0/17)
v2 corpus (42 / 37)	42/42 (0/37)	41/42 (0/37)	42/42 (0/37)
Sensitivity: gpt-oss-120b judge, blind corpus only			
v3 blind (21 / 24)	6/21 (3/24)	21/21 (5/24)	21/21 (7/24)

TABLE IX
AGENT IN THE LOOP, 80 INJECAGENT CASES PER COLUMN (40 DIRECT-HARM + 40 DATA-STEALING). COUNTS (95% CI). "EXECUTED" = EMITTED AND AUTHORIZED.

	Base	Enhanced
Undefended: attacker call emitted	8/80 [5.2, 18.5]	4/80 [2.0, 12.2]
Rules flagged the tool response	1/80	80/80
Executed, rules filter[a]	8/80 [5.2, 18.5]	0/80 [0.0, 4.6]
Executed, task-scoped grants (P1)	0/80 [0.0, 4.6]	0/80 [0.0, 4.6]
Executed, toolkit grants + taint gating (P3)	0/80 [0.0, 4.6]	0/80 [0.0, 4.6]
Executed, rules filter + P3	0/80 [0.0, 4.6]	0/80 [0.0, 4.6]

[a]Assumes a redacted tool response yields no attacker call; the agent was not re-run on redacted input.

role assignment), and none of the 37 and 17 benign messages of the other two corpora.

The sensitivity row is a warning about how far to trust any single false-positive figure. A larger judge reached the same detection but flagged five of the 24 benign blind messages on its own (seven in combination), because it treated some benign operational requests—reading a config, calling a health endpoint—as risky. The false-positive rate therefore depends strongly on the judge, and with 24 benign messages the intervals are wide. We did not calibrate the 0.5 and 0.6 cut-offs. The judge also adds a model call per message (not timed here) and is itself an injection target: escalate-only operation stops a fooled judge from *admitting* an attack but not from *blocking* a legitimate message.

C. An agent in the loop

We give the gpt-oss-20b agent a simplified single-turn setting in the spirit of InjecAgent: the user's instruction, the user's tool call already made, the tool response carrying the injection, and function definitions for the user's tool and the attacker's tool(s). We count an attack as *emitted* when the agent's next message calls an attacker tool, and then apply each defense to what the agent emitted (Table IX). With the injection removed from the tool response (a 40-case control), the agent never called an attacker tool (0 of 40; upper bound 8.8%), so the emitted calls are attributable to the injection.

The agent resisted most injections: it emitted the attacker's call in 8 of 80 base cases (10.0%) and, in only 4 of 80 (5.0%) when the override phrase was added (Fisher exact $p = 0.37$ between the two). This contrasts with InjecAgent's report that its hacking prompt nearly doubled the attack success of a ReAct-prompted GPT-4 [9]; we did not investigate why, and one untested possibility is that newer models are trained against this widely published phrase. The rules filter, which flags the tool response in 80 of 80 enhanced cases but only 1 of 80 base cases (and not one of the eight that succeeded), therefore removes every enhanced attempt and none of the base ones. Every emitted attacker call was denied under the least-privilege policies: 0 of 80 executed, against 8 of 80 undefended on base cases (Fisher exact $p = 0.0065$) and 4 of 80 on enhanced cases ($p = 0.12$, not significant at this sample size). We do not pool the two columns because they reuse the same 80 cases.

Reading this correctly. (i) One agent, one model family, and a low base rate: the power to detect a policy effect is limited, and the result that the policy executed nothing is partly structural, as in Sec. X, because the attacker's tools are almost never granted to the user's task. (ii) Our single-turn function-calling setup is not InjecAgent's ReAct protocol, so these rates are not comparable to published ones. (iii) For data-stealing cases we score only the first attacker call. (iv) We measure no utility cost.

TABLE X
ORACLE-GUIDED EVASION (93 SEEDS; 95% CI). "INITIALLY BLOCKED" = SEEDS THE DETECTOR BLOCKED BEFORE ANY REWRITING.

Detector	Initially blocked	Evaded after of those	Median queries
v1 keyword (τ=60)	5/93	5/5 = 100% [56.6, 100]	2
v2 rules (τ=60)	78/93	70/78 = 89.7%	68
v2, untrusted (τ=40)	78/93	69/78 = 88.5%	80

IX. RESULTS IV: AN ADAPTIVE ADVERSARY

We implement a tier-A2 attacker as an oracle-guided search: for each seed attack it queries the detector on rewrites drawn from five meaning-preserving operators (synonym substitution, filler insertion, intra-word hyphenation, adjacent-letter typos, and phrase translation to Spanish/French/German) and their depth-2 compositions (budget 1,500 compositions per seed, fixed RNG seed), stopping at the first rewrite scoring below the threshold. Seeds are the 93 pooled attacks.

Table X: v2 raises the attacker's cost by about $30\times$ in queries but does not stop them, as predicted. The eight (nine at $\tau = 40$) remaining seeds resist all text-level operators but evade through *decoder-dependent* encodings (base32, reversal, Caesar shift) that work only if the target model decodes them unprompted; we report them separately and do not count them as evasions in the headline numbers.

Caveats. (i) The operators are meaning-preserving for a human reader, but we did *not* verify that a target LLM still obeys the rewritten text (hyphenation in particular may degrade comprehension); "evades the detector" is not "attack succeeds". (ii) The library was written by the detector's authors. (iii) The search is score-guided black-box, not an LLM- or gradient-guided attacker, and so is a lower bound on adaptive strength. Importantly, none of these rewrites alters

TABLE XI
ATTACKER GOAL COMPLETED UNDER ASSUMED COMPROMISE (1,054 CASES). "ANY STEP" = AT LEAST ONE ATTACKER CALL AUTHORIZED.

Policy	Goal completed	Any step
P0 no policy	1054 (100%)	1054
P1 task-scoped grants	0 (0%)	1
P2 toolkit-scoped grants	1 (0.1%)	74
P3 toolkit-scoped + taint gating	0 (0%)	8
P4 task-scoped + taint gating	0 (0%)	1

which tools an attacker's call needs, which is why the next experiment is unaffected.

X. RESULTS V: LEAST-PRIVILEGE ENFORCEMENT UNDER ASSUMED COMPROMISE

Under tier H we assume the injection fully succeeds and the agent attempts the attacker's tool calls. The context of those calls is the tool response that carried the injection, hence tainted. We evaluate five policies on all 1,054 InjecAgent cases using the real `ToolPolicy` code:

- **P0**: no policy (every call executes).
- **P1**: *task-scoped* grants: only the user's own tool.
- **P2**: *toolkit-scoped* grants: every tool in the user tool's toolkit (the coarser permissions typical of integrations).
- **P3**: P2 plus taint gating (non-read calls from tainted context need human approval).
- **P4**: P1 plus taint gating.

Tool sensitivity is assigned by a verb-prefix heuristic fixed *before* running (`Get/Search/List/...`→read; `Delete/Cancel/...`→destructive; `Send/Share/Transfer/...`→external; `TerminalExecute`→destructive; otherwise write); results depend on it.

Table XI shows that, with the agent assumed fully hijacked, goals complete in 0 of 1,054 cases under P1, P3 and P4 and 1 under P2. Taint gating reduces partially authorized attacks from 74 (P2) to 8 (P3). The user's own tool call (trusted, untainted context) is authorized in all 1,054 cases under every policy.

Why this must be read carefully. The result is largely *structural*: only one of the benchmark's 63 distinct attacker tools is also a user tool, and just 4 of 1,054 cases have all attacker tools inside the user's own toolkit. Any sensible task-scoped grant therefore denies nearly everything; the interesting content is that coarser toolkit grants (P2) leak 74 partial attacks and that taint gating closes most of them. The result also *assumes* that the set of tools a task needs is known in advance (a planning step, as in CaMeL [14]); that attacker arguments are not inspected (we grant them the benefit of the doubt); and it does *not* measure the utility cost: we only verify the user's primary call, not legitimate multi-step workflows that read untrusted content and then act, which taint gating would hold for approval. Utility under real workflows is the main outstanding cost question.

TABLE XII
LATENCY IN MILLISECONDS (REFERENCE MACHINE, SEC. V).

Operation	n	Mean	Median	p95	Max
v1 scoring, 74 chars	20,000	0.0250	0.0189	0.0430	18.83
v1 scoring, 3 KB	2,000	0.0383	0.0366	0.0493	0.196
v2 scoring, 74 chars	5,000	0.899	0.692	1.93	8.06
v2 scoring, 1 KB	2,000	11.67	10.81	19.38	40.67
v2 scoring, 3 KB	1,000	28.98	26.26	45.76	111.8
v2 intercept + log, 74 chars	3,000	1.894	1.660	3.55	21.85
v2 intercept + log, 3 KB	1,000	42.17	40.49	72.52	139.4
Envelope sign	5,000	0.0350	0.0286	0.0576	0.900
Envelope verify (HMAC + replay)	5,000	0.0301	0.0232	0.0582	0.347
Tool authorize (arg. scan on)	5,000	1.222	1.107	2.19	43.23

XI. RESULTS VI: OVERHEAD AND RESOURCE EXHAUSTION

A. Latency

Table XII gives measured cost (Sec. V for method and hardware). v2 scoring (median 0.69 ms) is about 37× slower than v1 (0.019 ms) on short messages and grows about linearly with length, at roughly 9–11 ms per kilobyte, so a 3 KB message costs on the order of 26 ms. Signing, verification and tool authorization are cheap relative to scoring. All of this is small beside a model call (seconds), but it is not negligible, and our first manuscript's "17 μs" figure applies to v1 only.

B. Algorithmic-complexity exposure

Because a security filter sits in front of the model, a message that is expensive to scan is a denial-of-service vector against the pipeline. We fuzzed v2 with 18 pattern-dense and degenerate inputs (repeated trigger phrases, punctuation runs, spaced letters, long encoded-looking blobs, Unicode tag runs) at 10,000 and 100,000 characters. We found *no catastrophic backtracking*—cost stayed linear—but a 100,000-character message took up to 2.1 s and a 1 MB message would take on the order of 20 s. We therefore added the fail-closed bound $L_{\text{scan}} = 32{,}768$ characters (Sec. IV); it does not change any result in this paper, since every corpus message is far shorter. The alternative "scan head and tail only" mode is available but lets an attacker hide an injection in the unscanned middle, so it is not the default.

XII. DISCUSSION

What the rules layer is worth. It is an inexpensive, inspectable, auditable noise filter: it catches the unsophisticated and the obviously obfuscated (encoded 3/3, tag smuggling, chat-template spoofing, exfiltration URLs), it produces a readable reason for each block, and normalization demonstrably matters. It is not a security boundary: against paraphrase, other languages, plain-request injections and an adaptive attacker it fails, and we expect those failures to be the realistic case.

What a judge adds, and costs. The one place content inspection helped on plain-request indirect injection was an LLM judge: Qwen3.8-27B took InjecAgent detection from 2 to 90 of 100 and blind-corpus detection from 6 to 21

of 21, which neither the rules nor Llama Prompt Guard 2 classifier approached. The price is a model call per message, an additional injection target, and a false-positive rate that varied from 2 to 5 of 24 benign blind messages between the Qwen3.8-27B and the gpt-oss-120b judge. We would deploy such a judge only on untrusted or tainted content (its scope option), with its output feeding the audit log and an approval step rather than silently blocking.

What held. Deterministic controls that do not depend on understanding the text—message authentication, taint propagation, default-deny grants, argument validation, approval gating—held where detection did not, in the one evaluation where we could test them on external data. This matches the position of Beurer-Kellner et al. [13], CaMeL [14] and Fides [15] and the 4.2%-versus-31.2% result of Paul and Nandy [4], though our evaluation is far more limited than theirs (Sec. XIII).

Recommended deployment. Treat SecureFlow as three things: an audit/observability point; a first-pass filter that spends little and catches the cheap attacks; and a tool-authorization boundary whose grants are as narrow as the task allows, with human approval for non-read actions that follow untrusted input. The scoring threshold should be chosen by fixing a false-positive budget as in Liu et al. [7]; on our pooled data no threshold reaches both low BER and zero FPR (Fig. 2).

XIII. Threats to Validity, Reproducibility, and Planned Work

A. Limitations of the evidence

Limited end-to-end measurement. The detection, adaptive and policy experiments place no model in the loop. The one end-to-end experiment (Sec. VIII) uses a single agent (gpt-oss-20b), 80 cases per variant, a simplified single-turn protocol and only the first attacker call; the agent emitted attacks rarely, so statistical power is low and the enhanced-variant policy effect is not significant ($p = 0.12$). It makes no claim about the attack success rate of any other model or of a production pipeline. (The reference worker's response-keyword compliance check would mislabel a correct refusal as compliance and was not used.)

Corpus authorship. The v1–v3 corpora and the rules were written by the same authors. We mitigate with a blind corpus (v3) and an independent benchmark (InjecAgent), but v3 is still author-written and small (21/24), its interval wide (BER 50.0–86.2%), and the English-only composition understates multilingual exposure.

Tuning contamination. v1 and v2 corpora informed the v2 rules and are reported only as such. The v3 and InjecAgent results are the generalization evidence; rules were not changed after InjecAgent was first run.

Small benign sets. 81 pooled benign messages and 17 InjecAgent templates give wide FPR intervals (pooled 1.3–10.3%; InjecAgent upper bound 18.4%). The FPR on realistic security-operations traffic is unmeasured.

Policy evaluation scope. Results depend on task-scoped grants being known, on the sensitivity heuristic, and on trusted/tainted labelling being applied correctly by the application; utility cost is unmeasured (Sec. X).

Adaptive attacker. Operators are author-written and meaning preservation for an LLM is unverified (Sec. IX).

Semantic judge and classifier. Each was evaluated with one model per role on small samples (24, 17 and 37 benign messages; 100 InjecAgent attacks). The judge's false-positive rate moved from 2 to 5 of 24 with the model chosen, cut-offs were not calibrated, its latency and cost were not measured, and we did not attack the judge itself. The classifier (Llama Prompt Guard 2) is a small 86-million-parameter model.

B. Limitations of the mechanism

Signing uses shared secrets within one trust domain; cross-domain use needs asymmetric signatures. Taint is a single untrusted bit that is only as accurate as the application's labelling; SecureFlow cannot see data flows it is not told about. The hash chain detects tampering but does not prevent it. Rules are English-centric. Suppression does not constrain tools the application does not route through `authorize`.

C. Reproducibility

Everything is public. Code: https://github.com/piyushrajyadav/Secure-Flow. The released test data are under `data/corpora/` (`v1.jsonl`, `v2.jsonl`, `v3.jsonl`; SHA-256 checksums in `MANIFEST.json`); all result files are under `data/results/`, each stamped with the producing git commit; InjecAgent is fetched at a pinned commit by a script rather than redistributed. `python -m experiments.run_all` regenerates every number here in about seven minutes without a model or GPU; the 75-test suite (`python -m pytest tests`) includes regression guards on the corpora. Detection, ablation, InjecAgent and adaptive results are deterministic (the adaptive search uses a fixed RNG seed); latency varies with hardware and load. The model-in-the-loop results (Sec. VIII) are re-derived from the released request–response cache (`data/results/llm_cache.jsonl`, pruned to exactly the entries the experiments use) by running `python -m experiments.run_all` with replay enabled, which needs no network access and no credentials; we verified that replay reproduces every count reported here. Re-running against a live service needs a key and may differ slightly because hosted-model outputs can drift.

D. Ethics and responsible use

The attack strings are generic prompt-injection patterns and public-benchmark text; we disclose no vulnerability in any third-party product. The adaptive-rewrite library defeats a detector we also release; we publish it because the detector is open source and a tier-A2 adversary can build the same, and because honest evaluation requires it. The released corpora contain adversarial text and should be treated as test data, never as instructions.

E. Planned work

None of the following has been performed, and none is reported as a result.

1) **Larger-scale end-to-end evaluation**: several agent models and families, InjecAgent's own ReAct protocol and metric, AgentDojo [8], repetitions per case, and per-case success criteria with inter-annotator agreement. The agent experiment of Sec. VIII is a small first step, not a substitute.
2) **Judge calibration and attack**: choose the cut-offs under a fixed false-positive budget [7], measure latency and cost, compare several judge models and a larger safeguard model [11], and attack the judge with adaptive injections.
3) **A larger, independently authored, multilingual benign and attack corpus**, including content-moderation prompts that legitimately quote injections.
4) **Utility measurement of taint gating** on realistic multi-step workflows.
5) **LLM-guided adaptive attacks** and verification that rewrites remain effective against target models.

XIV. Conclusion

We set out to learn what a transparent checkpoint on the inter-agent edge can and cannot do. Its content-scoring half does less than it appears to: v1 forwarded 94.6% of 93 attacks because of how its weights and threshold interact; v2 cuts that to 16.1% on data it was tuned near, but forwards 71.4% of a blind corpus and 98.4% of plain-request injections on an external benchmark, as does the 86-million-parameter Prompt Guard 2 classifier, and, once an oracle-guided attacker rewrites them, about 90% of the attacks it initially detects. A Qwen3.8-27B judge recovered most of the missing recall (90 of 100 on that benchmark; 21 of 21 on the blind corpus) at a false-positive cost that depended strongly on which judge was used. The enforcement half did the most to limit damage: with a gpt-oss-20b agent in the loop, every emitted attacker call was denied by least-privilege authorization (0 of 80 executed in each variant, against 8 of 80 undefended on the base cases), and assuming total hijack the attacker's goal completed in none of 1,054 external cases—though both results are partly structural and the utility cost is unmeasured. We offer SecureFlow as an audit point, a cheap first filter, an optional judge for untrusted content, and a tool-authorization boundary, to be deployed beneath an isolation design rather than in place of one. We would rather report a defense whose detector does not generalize than a number that does not mean what it says.

Acknowledgment

The authors thank the Department of Computer Science and Engineering, Institute of Engineering and Management, Kolkata, for its support, and the open-source communities behind LangGraph, LangChain, and Ollama [24].

References

[1] LangChain AI, "LangGraph documentation," 2024. [Online]. Available: https://langchain-ai.github.io/langgraph/. [Accessed: Oct. 16, 2026].
[2] K. Greshake, S. Abdelnabi, S. Mishra, C. Endres, T. Holz, and M. Fritz, "Not what you've signed up for: Compromising real-world LLM-integrated applications with indirect prompt injection," *arXiv preprint arXiv:2302.12173*, 2023.
[3] D. Lee and M. Tiwari, "Prompt infection: LLM-to-LLM prompt injection within multi-agent systems," *arXiv preprint arXiv:2410.07283*, 2024.
[4] R. K. Paul and S. Nandy, "Beyond single-model injection: A threat model and defense architecture for prompt injection in multi-agent systems," *arXiv preprint arXiv:2609.22949*, Sep. 2026.
[5] OWASP GenAI Security Project, "OWASP top 10 for agentic applications for 2026," Dec. 2025. [Online]. Available: https://genai.owasp.org/resource/owasp-top-10-for-agentic-applications-for-2026/.
[6] F. Perez and I. Ribeiro, "Ignore previous prompt: Attack techniques for language models," in *Proc. NeurIPS Workshop on ML Safety*, 2022. [Online]. Available: https://arxiv.org/abs/2211.09527.
[7] Y. Liu, Y. Jia, R. Geng, J. Jia, and N. Z. Gong, "Formalizing and benchmarking prompt injection attacks and defenses," in *Proc. 33rd USENIX Security Symp. (USENIX Security 24)*, Philadelphia, PA, USA, Aug. 2024, pp. 1831–1847.
[8] E. Debenedetti, J. Zhang, M. Balunović, L. Beurer-Kellner, M. Fischer, and F. Tramèr, "AgentDojo: A dynamic environment to evaluate prompt injection attacks and defenses for LLM agents," in *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*, 2024; *arXiv preprint arXiv:2406.13352*.
[9] Q. Zhan, Z. Liang, Z. Ying, and D. Kang, "InjecAgent: Benchmarking indirect prompt injections in tool-integrated large language model agents," in *Findings of the Association for Computational Linguistics: ACL 2024*, 2024.
[10] T. Rebedea, R. Dinu, M. N. Sreedhar, C. Parisien, and J. Cohen, "NeMo Guardrails: A toolkit for controllable and safe LLM applications with programmable rails," in *Proc. 2023 Conf. Empirical Methods in Natural Language Processing: System Demonstrations*, Singapore, Dec. 2023, pp. 431–445. doi: 10.18653/v1/2023.emnlp-demo.40.
[11] H. Inan *et al.*, "Llama Guard: LLM-based input-output safeguard for human-AI conversations," *arXiv preprint arXiv:2312.06674*, 2023.
[12] K. Hines, G. Lopez, M. Hall, F. Zarfati, Y. Zunger, and E. Kiciman, "Defending against indirect prompt injection attacks with spotlighting," *arXiv preprint arXiv:2403.14720*, 2024.
[13] L. Beurer-Kellner *et al.*, "Design patterns for securing LLM agents against prompt injections," *arXiv preprint arXiv:2506.08837*, 2025.
[14] E. Debenedetti *et al.*, "Defeating prompt injections by design," *arXiv preprint arXiv:2503.18813*, 2025.
[15] M. Costa, B. Köpf, A. Kolluri, A. Paverd, M. Russinovich, A. Salem, S. Tople, L. Wutschitz, and S. Zanella-Béguelin, "Securing AI agents with information-flow control," *arXiv preprint arXiv:2505.23643*, 2025.
[16] OpenAI, "gpt-oss-120b & gpt-oss-20b model card," *arXiv preprint arXiv:2508.10925*, Aug. 2025.
[17] Meta, "Llama Prompt Guard 2 (22M and 86M) model card," 2025. [Online]. Available: https://www.llama.com/docs/model-cards-and-prompt-formats/prompt-guard/.
[18] Qwen Team, Alibaba Group, "Qwen3.8-27B," model card (open weights, Apache-2.0), Aug. 2026. [Online]. Available: https://huggingface.co/Qwen/Qwen3.8-27B.
[19] H. Krawczyk, M. Bellare, and R. Canetti, "HMAC: Keyed-hashing for message authentication," IETF, RFC 2104, Feb. 1997.
[20] E. B. Wilson, "Probable inference, the law of succession, and statistical inference," *J. Amer. Statist. Assoc.*, vol. 22, no. 158, pp. 209–212, 1927.
[21] International Organization for Standardization, "ISO/IEC 23894: Information technology — Artificial intelligence — Guidance on risk management," 2023.
[22] National Institute of Standards and Technology, "Artificial intelligence risk management framework (AI RMF 1.0)," NIST AI 100-1, Jan. 2023.
[23] OWASP Foundation, "OWASP top 10 for large language model applications," 2025. [Online]. Available: https://owasp.org/www-project-top-10-for-large-language-model-applications/.
[24] Ollama Project, "Ollama: Local large language models," 2024. [Online]. Available: https://ollama.com/. [Accessed: Oct. 16, 2026].